

Live Session, Feb. 15

Quicksort and Closest Pair of Points

CS 4102: Algorithms

Spring 2021

Mark Floryan and Tom Horton

Monday, Feb. 15

- ▶ Gradescope is live (link via Collab to sign up)
 - ▶ Made some cosmetic changes to programming hws
- ▶ Recommended schedule up on schedule page of website.
 - ▶ Recommended that “distancing” be done this week
- ▶ Sample HW assignment is up on webpage and GS
 - ▶ For your benefit if you want it, not required.
- ▶ Office hours going ok?
- ▶ Don’t forget Wednesday is a break day so no class!
- ▶ Today we will do quicksort and closest pair of points!
- ▶ Correction: Closest pair of points readings: CLRS 33.4
- ▶ Self-assessment “quiz”: <https://forms.gle/UY83NtmQT2PBdkXT8>

Quicksort and Partition

Readings: CLRS Chapter 7 (not 7.4.2)

Quicksort's Strategy

- ▶ Called on subsection of array from *first* to *last*
 - ▶ Like mergesort
- ▶ First, choose some element in the array to be the ***pivot*** element
 - ▶ Any element! Doesn't matter for correctness.
 - ▶ Often the first item. For us, the last. Or, we often move some element into the last position (to get better efficiency)
- ▶ Second, call ***partition***, which does two things:
 - ▶ Puts the pivot in its proper place, i.e. where it will be in the correctly sorted sequence
 - ▶ All elements below the pivot are less-than the pivot, and all elements above the pivot are greater-than
- ▶ Third, use quicksort recursively on both sub-lists

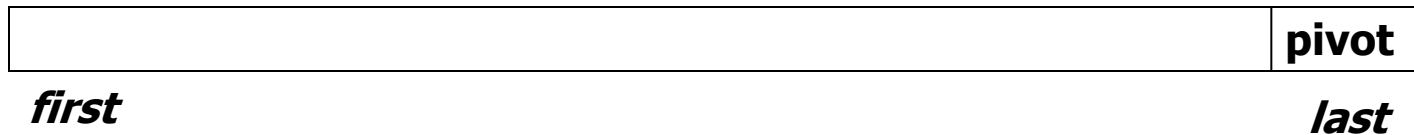
Quicksort is Divide and Conquer

- ▶ **Divide:** select **pivot** element p , **Partition**(p)
- ▶ **Conquer:** recursively sort left and right sublists
- ▶ **Combine:** Nothing!

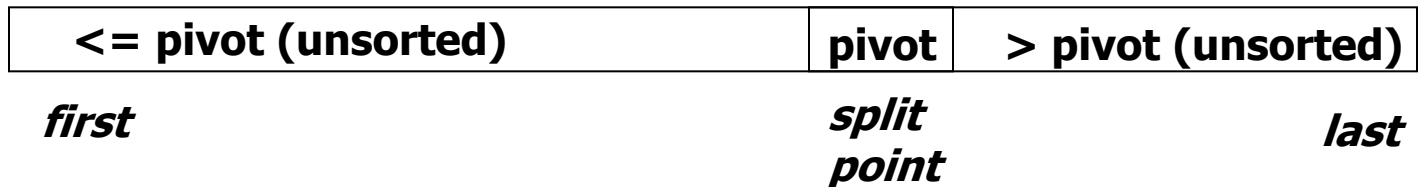
Contrast to mergesort,
where divide is simple and combine is work

Quicksort's Strategy (a picture)

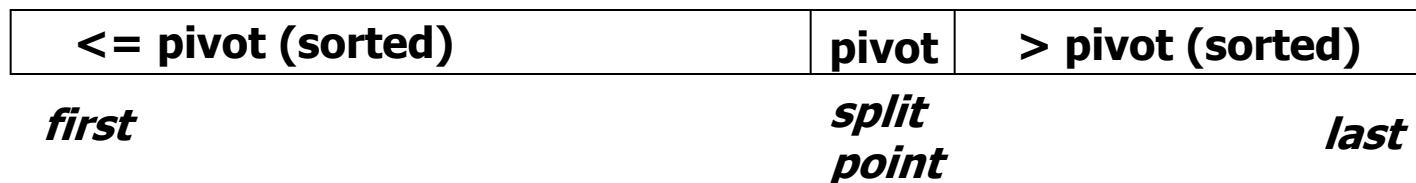
- ▶ Use last element as pivot (or pick one and move it there)



- ▶ After call to partition...



- ▶ Now sort two parts recursively and we're done!



- ▶ Note that splitPoint may be anywhere in *first..last*
- ▶ Note our assumption that all keys are distinct

Quicksort Code


Input Parameters: *list, first, last*

Output Parameters: *list*

```
def quicksort(list, first, last):  
    if first < last:  
        q = partition(list, first, last)  
        quicksort(list, first, q-1)  
        quicksort(list, q+1, last)  
    return
```

Strategy for Lomuto's Partition

- Invariant: At any point:
 - i indexes the right-most element $\leq \text{pivot}$
 - $j-1$ indexes the right-most element $> \text{pivot}$



- ▶ Strategy:
 - ▶ Look at next item $a[j]$
 - ▶ If that item $> \text{pivot}$, all is well!
 - ▶ If that item $< \text{pivot}$, increment i and then swap items at positions i and j
 - ▶ When done, swap pivot with item at position $i+1$
- ▶ Number of comparisons: $n-1$

Lomuto's Partition: Code

Input Parameters: *list, first, last*


Output Parameters: *list*.

Return value: the split point

```
def partition(list, first, last):  
    pval = list[last]  
    i = first-1  
    for j in range(first, last): # first up to before last  
        if list[j] <= pval:  
            i = i+1  
            (list[i], list[j]) = (list[j], list[i]) # swap!  
    (list[last], list[i+1]) = (list[i+1], list[last]) # swap!  
    return i+1
```

Partition this: [c t o a m b f]

Partition this: [c t o a m b f]



- Strategy:
 - Look at next item $a[j]$
 - If that item $>$ pivot, all is well!
 - If that item $<$ pivot, increment i and then swap items at positions i and j
 - When done, swap pivot with item at position $i+1$

Partition this: [c t o a m b f]

- ▶ Was that a “good” result? Correct? Desirable for some other reason?
- ▶ Let’s do a randomized partition now!

Efficiency of Quicksort

- ▶ Partition divides into two sub-lists, perhaps unequal size
 - ▶ Depends on value of pivot element
- ▶ Recurrence for Quicksort
$$T(n) = \text{partition-cost} + T(\text{size of 1st section}) + T(\text{size of 2nd section})$$
- ▶ If divides equally, $T(n) = 2 T(n/2) + n-1$
 - ▶ Just like mergesort
 - ▶ Solve by substitution or master theorem
$$T(n) \in \Theta(n \lg n)$$
- ▶ This is the best-case. But...

Worst Case of Quicksort

- ▶ What if divides in most unequal fashion possible?
 - ▶ One subsection has size 0, other has size $n-1$
 - ▶ $T(n) = T(0) + T(n-1) + n-1$
 - ▶ What if this happens every time we call partition recursively?

$$W(n) = \sum_{k=2}^n (k-1) \in \Theta(n^2)$$

- ▶ Uh oh. Same as insertion sort.
 - ▶ “Sorry Prof. Hoare – we have to take back that Turing Award now!”

Quicksort's Average Case

- ▶ Good if it divides equally, bad if most unequal.
 - ▶ Remember: when subproblems size 0 and $n-1$
 - ▶ Can worst-case happen?
Sure! Many cases. One is when elements already sorted. Last element is max, pivot around that. Next pivot is 2nd max...
- ▶ What's the average?
 - ▶ Much closer to the best case
 - ▶ A bad-split then a good-split is closer to best-case (pp. 176-178)
 - ▶ To prove $A(n)$, fun with recurrences!
 - ▶ The result: If all permutations are equal, then
$$A(n) \cong 1.386 n \lg n \text{ (for large } n)$$
- ▶ So very fast on average.
- ▶ And, we can take simple steps to avoid the worst case!

Avoiding Quicksort's Worst Case

- ▶ Make sure we don't pivot around max or min
 - ▶ Find a better choice and swap it with last element
 - ▶ Then partition as before
- ▶ Recall we get best case if divides equally
 - ▶ Could find median. But this costs $\Theta(n)$. Instead...
 - ▶ Choose a **random element** between first and last and swap it with the last element
 - ▶ Or, estimate the median by using the “median-of-three” method
 - ▶ Pick 3 elements (say, first, middle and last)
 - ▶ Choose median of these and swap with last. (Cost?)
 - ▶ If sorted, then this chooses real median. Best case!

Tuning Quicksort's Performance

- ▶ In practice quicksort runs fast
 - ▶ $A(n)$ is log-linear, and the “constants” are smaller than mergesort and heapsort
 - ▶ Often used in software libraries
 - ▶ So worth tuning it to squeeze the most out of it
 - ▶ Always do something to avoid worst-case
- ▶ Sort small sub-lists with (say) insertion sort
 - ▶ For small inputs, insertion sort is fine
 - ▶ No recursion, function calls
 - ▶ Variation: don't sort small sections at all.
After quicksort is done, sort entire array with **insertion sort**
 - ▶ It's efficient on almost-sorted arrays!

Quicksort's Space Complexity

- ▶ Looks like it's in-place, but there's a *recursion stack*
 - ▶ Depends on your definition: some people define *in-place* to not include stack space used by recursion
 - ▶ E.g. our CLRS algorithms textbook
 - ▶ Other books and people do “count” this
 - ▶ How much goes on the stack?
 - ▶ If most uneven splits, then $\Theta(n)$.
 - ▶ If splits evenly every time, then $\Theta(\lg n)$.
- ▶ Ways to reduce stack-space used due to recursion
 - ▶ Various books cover the details (not ours, though)
 - ▶ First, remove 2nd recursive call (tail-recursion)
 - ▶ Second, always do recursive call on smaller section

Summary: Quicksort

- ▶ Divide and conquer where divide does the heavy-lifting
- ▶ In worst-case, efficiency is $\Theta(n^2)$
 - ▶ But it's practical to avoid the worst-case
- ▶ On average, efficiency is $\Theta(n \lg n)$
- ▶ Better space-complexity than mergesort.
- ▶ In practice, runs fast and widely used
 - ▶ Many ways to tune its performance
- ▶ Various strategies for Partition
 - ▶ Some work better if duplicate keys
- ▶ More details? See Sedgewick's algorithms textbook
 - ▶ He's the expert! PhD on this under Donald Knuth

Closest Pair of Points

Readings: CLRS 33.4


Closest Pair of Points in 2D Space

Given:

A list of points

Return:

Distance of the pair of points that are closest together
(or possibly the pair too)



Closest Pair of Points: Naïve

Given:

A list of points

Return:


Distance of the closest pair of points

Naive Algorithm: $O(n^2)$

Test every pair of points,
return the closest.

We can do better!

$\Theta(n \log n)$




Closest Pair of Points: D&C

Divide: How?

At median x coordinate

Conquer:

Combine:



Closest Pair of Points: D&C


Divide:

At median x coordinate

Conquer:

Recursively find closest pairs from Left and Right

Combine:



Closest Pair of Points: D&C

Divide:


At median x coordinate

Conquer:

Recursively find closest pairs from Left and Right

Combine:

Return min of Left and Right pairs
Problem
?



Closest Pair of Points: D&C


Combine:

2 Cases:

1. Closest Pair is completely in Left or Right

2. Closest Pair Spans our “Cut”

Need to test points across the cut



Spanning the Cut

Define “runway” or
“strip” along the cut.


Combine:

2. Closest Pair Spanned our “Cut”

Need to test points
across the cut.

Bad approach: Compare
all points within $\delta =$
 $\min\{\delta_L, \delta_R\}$ of the cut.

How many are there?



Spanning the Cut

Define “runway” or
“strip” along the cut.

Combine:


2. Closest Pair Spanned our “Cut”

Need to test points
across the cut

Bad approach: Compare
all points within $\delta =$
 $\min\{\delta_L, \delta_R\}$ of the cut.

How many are there?

$$\begin{aligned} T(n) &= 2T\left(\frac{n}{2}\right) + \left(\frac{n}{2}\right)^2 \\ &= \Theta(n^2) \end{aligned}$$



Spanning the Cut


Combine:

2. Closest Pair Spanned our “Cut”

Need to test points
across the cut

We don't need to test all
pairs!

Don't need to test any
points that are $> \delta$ from
one another



Spanning the Cut

Combine:

2. Closest Pair Spanned our “Cut”


Consider points in strip in increasing y-order.

For a given point p , we can *prove* the 8th point and beyond is more than δ from p .

(pp. 1041-2 in CLRS)

So for each point in strip, check next 7 points in y-order.

$\Theta(n)$ **Better!**




Reducing Search Space

Combine:

Need to test points across the cut

Claim #1: if two points are the closest pair that cross the cut, then you can surround them in a box that's $2 \cdot \delta$ wide by δ tall.

Let's draw some examples.



Reducing Search Space


Assume you're checking in increasing y-order, and you've reached the first point of the closest pair.

Do you have to look at **all points above it** to be guaranteed to find the other point and the minimum distance?

No!

- Imagine you drew a box with its bottom at point's y-coordinate.
- See Claim #1.
- Claim #2: only 8 points can be in the box.

Claim #1: if two points are the closest pair that cross the cut, then you can surround them in a box that's $2 \cdot \delta$ wide by δ tall.



Reducing Search Space


Assume you're checking in increasing y-order, and you've reached the first point of the closest pair.

Do you have to look at **all points above it** to be guaranteed to find the other point and the minimum distance?

No!

- Imagine you drew a box with its bottom at point's y-coordinate.
- See Claim #1.
- Claim #2: only 8 points can be in the box.

Claim #1: if two points are the closest pair that cross the cut, then you can surround them in a box that's $2 \cdot \delta$ wide by δ tall.



Spanning the Cut

Combine:

2. Closest Pair Spanned our “Cut”


Consider points in strip in increasing y-order.

For a given point p , we can *prove* the 8th point and beyond is more than δ from p .

(pp. 1041-2 in CLRS)

So for each point in strip, check next 7 points in y-order.

$\Theta(n)$ **Better!**



Closest Pair of Points: Divide and Conquer

Initialization: Sort points by x -coordinate
(Later we'll also need to process points by y -coordinate, too.)


Divide: Partition points into two lists of points based on x -coordinate (split at the median x)

Conquer: Recursively compute the closest pair of points in each list

Base case?

Combine:

- Consider only points in the runway (x -coordinate within distance δ of median)
- Process runway points by y -coordinate
- Compare each point in runway to 7 points above it and save the closest pair
- Output closest pair among **left**, **right**, and **runway** points



Closest Pair of Points: Divide and Conquer

What is the running time?

$$\Theta(n \log n)$$

$$T(n)$$

$$T(n) = 2T(n/2) + \Theta(n)$$

Case 2 of Master's Theorem

$$T(n) = \Theta(n \log n)$$

$$\Theta(n \log n)$$

$$\Theta(1)$$

$$2T(n/2)$$

$$\Theta(n)$$

$$\Theta(1)$$

Initialization: Sort points by x -coordinate

Divide: Partition points into two lists of points based on x -coordinate (split at the median x)

Conquer: Recursively compute the closest pair of points in each list

Combine:

- Process runway points by y -coordinate and Compare each point in runway to 7 points above it and save the closest pair
- Output closest pair among **left**, **right**, and **runway** points

Summary for Closest Pair of Points

- ▶ Comparing all pairs is a brute-force fail
 - ▶ Except for small inputs
- ▶ Divide and conquer a big improvement
- ▶ Needed to find an efficient way for part of the combine step
 - ▶ Geometry came through for us here!
 - ▶ Only needed to look at constant number of points for each point in the strip
- ▶ Implementation subtleties
 - ▶ Don't want to sort the strip by y-coordinate in each recursive call
 - ▶ In initialization, create an “index” that lets you process all points in order by y-coordinate
 - ▶ (There are other ways to address this.)